
Data-Driven Stock Investment Decision Support

A Reinforcement Learning Approach to
Risk-Sensitive Sequential Portfolio Decisions

Master Thesis

Jannik Busse Guldmand

SUPERVISOR

Professor Melih Kandemir

SECOND READER

Professor Unknown



University of
Southern Denmark

June 1, 2026

Disclaimer

IN THIS CHAPTER Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nullam aliquam sem lacus, sodales gravida lacus volutpat feugiat. Sed aliquet eleifend sapien, ut lacinia magna porta vulputate. Curabitur in leo ac lacus volutpat vehicula at ac nisi. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Sed scelerisque lacinia neque, et imperdiet risus bibendum at. Nullam vitae arcu congue, lacinia dui vitae, convallis ipsum. Morbi auctor elit sit amet nibh cursus vulputate. Praesent faucibus congue risus, non elementum velit suscipit vel.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nullam aliquam sem lacus, sodales gravida lacus volutpat feugiat. Sed aliquet eleifend sapien, ut lacinia magna porta vulputate. Curabitur in leo ac lacus volutpat vehicula at ac nisi. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Sed scelerisque lacinia neque, et imperdiet risus bibendum at. Nullam vitae arcu congue, lacinia dui vitae, convallis ipsum. Morbi auctor elit sit amet nibh cursus vulputate. Praesent faucibus congue risus, non elementum velit suscipit vel.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nullam aliquam sem lacus, sodales gravida lacus volutpat feugiat. Sed aliquet eleifend sapien, ut lacinia magna porta vulputate. Curabitur in leo ac lacus volutpat vehicula at ac nisi. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Sed scelerisque lacinia neque, et imperdiet risus bibendum at. Nullam vitae arcu congue, lacinia dui vitae, convallis ipsum. Morbi auctor elit sit amet nibh cursus vulputate. Praesent faucibus congue risus, non elementum velit suscipit vel.

Preface

IN THIS CHAPTER Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nullam aliquam sem lacus, sodales gravida lacus volutpat feugiat. Sed aliquet eleifend sapien, ut lacinia magna porta vulputate. Curabitur in leo ac lacus volutpat vehicula at ac nisi. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Sed scelerisque lacinia neque, et imperdiet risus bibendum at. Nullam vitae arcu congue, lacinia dui vitae, convallis ipsum. Morbi auctor elit sit amet nibh cursus vulputate. Praesent faucibus congue risus, non elementum velit suscipit vel.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nullam aliquam sem lacus, sodales gravida lacus volutpat feugiat. Sed aliquet eleifend sapien, ut lacinia magna porta vulputate. Curabitur in leo ac lacus volutpat vehicula at ac nisi. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Sed scelerisque lacinia neque, et imperdiet risus bibendum at. Nullam vitae arcu congue, lacinia dui vitae, convallis ipsum. Morbi auctor elit sit amet nibh cursus vulputate. Praesent faucibus congue risus, non elementum velit suscipit vel.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nullam aliquam sem lacus, sodales gravida lacus volutpat feugiat. Sed aliquet eleifend sapien, ut lacinia magna porta vulputate. Curabitur in leo ac lacus volutpat vehicula at ac nisi. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Sed scelerisque lacinia neque, et imperdiet risus bibendum at. Nullam vitae arcu congue, lacinia dui vitae, convallis ipsum. Morbi auctor elit sit amet nibh cursus vulputate. Praesent faucibus congue risus, non elementum velit suscipit vel.

Abstract

This is temporary abstract text. The purpose of this dummy document is to create enough structure to evaluate the table of contents, chapter ordering, section hierarchy, and later typography choices. The content is placeholder text and should be replaced by the final thesis abstract.

Contents

Disclaimer	i
Preface	ii
Abstract	iii
Contents	iv
1 Introduction	1
1.1 Box gallery (visual test)	1
1.2 Motivation	1
1.3 Problem Formulation	1
1.4 Research Questions	2
1.5 Scope and Delimitations	2
1.6 Thesis Structure	2
2 Background	3
2.1 Financial markets	3
2.2 The Investor	3
2.3 Portfolio theory	3
2.4 Risk, return, and uncertainty	3
2.5 Inflation and macroeconomic context	3
2.6 Decision support systems in finance	3
2.7 Reinforcement learning for investment decisions	4
3 Related Work	5
3.1 Direct Predecessors	5
3.2 Closest Application Work	5
4 Literature Survey	6
4.1 State of the Art	6
4.2 Reinforcement Learning for Algorithmic Trading and Portfolio Management	6
4.3 Distributional Reinforcement Learning	6
4.4 Uncertainty- and Risk-Sensitive Reinforcement Learning	6

4.5	Positioning: The Research Gap	6
4.6	Chosen Solution	6
4.7	Alternative Solutions	6
5	Theory	7
5.1	Markov Decision Processes	7
5.2	Dynamic Programming	7
5.3	Q-Learning	7
5.4	Deep Q-Networks (DQN)	7
5.5	The Limitation of a Single Scalar Value	7
5.6	Distributional Reinforcement Learning	7
5.7	Quantile Regression and the Quantile Huber Loss	7
5.8	Implicit Quantile Networks (IQN)	7
5.9	Risk-Sensitive Policies via Distortion Measures	7
5.10	Evidential Deep Learning	7
5.11	The Complete Chain	7
6	Methodology	8
6.1	Research Design	8
6.2	Problem Formulation as an MDP	8
6.3	D-IQN-DSS: The Decision Engine	8
6.4	Hierarchical Decision Policy (HDP)	8
6.5	Epistemic Uncertainty via EDL	9
6.6	The Fusion Rule	9
6.7	Risk-Sensitive Decision Scoring	9
6.8	Scope and Limitations	9
6.9	Experimental Design	9
6.10	Evaluation Protocol	9
	Acknowledgements	10
A	Additional Material	11
A.1	Experiment Configuration	11
A.2	Additional Tables	11
A.3	Box gallery	11

Introduction

1.1 Box gallery (visual test)

Definition 1.1.1 (Invariants of unimodular symmetric forms)

Let Z be a finitely-generated free abelian group, and let $q : Z \times Z \rightarrow \mathbb{Z}$ be a unimodular form. To it we assign rank, type (parity), and signature.

1.2 Motivation

Stock investors make sequential decisions under uncertainty. The motivation for this thesis is to explore how reinforcement learning can support such decisions in a transparent and risk-aware way.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer vitae nisl sed neque facilisis tincidunt. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia curae. This dummy text is included only to create realistic page flow and table of contents entries while the LaTeX template is being developed.

Praesent non sem vel lorem varius hendrerit. Donec at massa sit amet ipsum tristique convallis. The final thesis text should replace this placeholder once the structure and visual style have been stabilized.

1.3 Problem Formulation

The central problem is how to design a stock investment decision support system that provides meaningful recommendations while avoiding look-ahead bias.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer vitae nisl sed neque facilisis tincidunt. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia curae. This dummy text is included only to create realistic page flow and table of contents entries while the LaTeX template is being developed.

Praesent non sem vel lorem varius hendrerit. Donec at massa sit amet ipsum tristique convallis. The final thesis text should replace this placeholder once the structure and visual style have been stabilized.

1.4 Research Questions

The thesis investigates how reinforcement learning, distributional value estimation, and point-in-time data handling can be combined in a decision support system.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer vitae nisl sed neque facilisis tincidunt. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia curae. This dummy text is included only to create realistic page flow and table of contents entries while the LaTeX template is being developed.

Praesent non sem vel lorem varius hendrerit. Donec at massa sit amet ipsum tristique convallis. The final thesis text should replace this placeholder once the structure and visual style have been stabilized.

1.5 Scope and Delimitations

The thesis focuses on a proof of concept rather than a production trading platform.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer vitae nisl sed neque facilisis tincidunt. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia curae. This dummy text is included only to create realistic page flow and table of contents entries while the LaTeX template is being developed.

Praesent non sem vel lorem varius hendrerit. Donec at massa sit amet ipsum tristique convallis. The final thesis text should replace this placeholder once the structure and visual style have been stabilized.

1.6 Thesis Structure

The thesis is organized into background, related work, methodology, theory, experiments, proof of concept, results, conclusion, and future perspectives.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer vitae nisl sed neque facilisis tincidunt. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia curae. This dummy text is included only to create realistic page flow and table of contents entries while the LaTeX template is being developed.

Praesent non sem vel lorem varius hendrerit. Donec at massa sit amet ipsum tristique convallis. The final thesis text should replace this placeholder once the structure and visual style have been stabilized.

Background

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.1 Financial markets

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.2 The Investor

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.3 Portfolio theory

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.4 Risk, return, and uncertainty

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.5 Inflation and macroeconomic context

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.6 Decision support systems in finance

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

2.7 Reinforcement learning for investment decisions

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

Related Work

3.1 Direct Predecessors

3.2 Closest Application Work

Literature Survey

- 4.1 State of the Art**
- 4.2 Reinforcement Learning for Algorithmic Trading and Portfolio Management**
- 4.3 Distributional Reinforcement Learning**
- 4.4 Uncertainty- and Risk-Sensitive Reinforcement Learning**
- 4.5 Positioning: The Research Gap**
- 4.6 Chosen Solution**
- 4.7 Alternative Solutions**

Theory

- 5.1 Markov Decision Processes
- 5.2 Dynamic Programming
- 5.3 Q-Learning
- 5.4 Deep Q-Networks (DQN)
- 5.5 The Limitation of a Single Scalar Value
- 5.6 Distributional Reinforcement Learning
- 5.7 Quantile Regression and the Quantile Huber Loss
- 5.8 Implicit Quantile Networks (IQN)
- 5.9 Risk-Sensitive Policies via Distortion Measures
- 5.10 Evidential Deep Learning
- 5.11 The Complete Chain

Methodology

This chapter describes how the general theory is shaped into the concrete system and how it is evaluated.

6.1 Research Design

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.2 Problem Formulation as an MDP

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

State space

The state $s_t \in \mathcal{S}$ encodes the information available to the agent at decision time t . It comprises the normalized price window, technical indicators, and the current portfolio allocation, so that $s_t = (\mathbf{p}_{t-w:t}, \mathbf{x}_{t-1})$ where w is the lookback length.

Action space

The action $a_t \in \mathcal{A}$ is the discrete trading decision applied to each asset, drawn from $\{\text{SELL}, \text{HOLD}, \text{BUY}\}$ at fixed position sizes.

Reward design

The reward r_t is the realized log-return of the portfolio over the holding period, optionally adjusted by a risk penalty.

6.3 D-IQN-DSS: The Decision Engine

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.4 Hierarchical Decision Policy (HDP)

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.5 Epistemic Uncertainty via EDL

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.6 The Fusion Rule

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.7 Risk-Sensitive Decision Scoring

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.8 Scope and Limitations

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.9 Experimental Design

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

6.10 Evaluation Protocol

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

Point-in-time walk-forward backtesting.

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

Metrics.

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi ar

Acknowledgements

Acknowledgements

First of all, I would like to thank the mathematical community whose documents make LaTeX typography worth studying. This section is included to test the visual spacing of front matter pages.

Additional Material

A.1 Experiment Configuration

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer facilisis, erat vel porttitor volutpat, neque erat dignissim ipsum, vitae commodo ligula turpis sed nulla.

A.2 Additional Tables

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Integer facilisis, erat vel porttitor volutpat, neque erat dignissim ipsum, vitae commodo ligula turpis sed nulla.

A.3 Box gallery

Definition A.3.1 (Risk-sensitive policy)

A policy π_β is risk-sensitive if it maximizes a distorted expectation $Q_\beta(x, a)$ rather than the plain expectation $\mathbb{E}[Z^\pi(x, a)]$.

Theorem A.3.2 (Bellman optimality)

The optimal action-value function Q^* is the unique fixed point of the Bellman optimality operator $\mathcal{T}$, satisfying $Q^* = \mathcal{T}Q^*$.

Proposition A.3.3 (Monotonicity of the distortion)

If β is a non-decreasing distortion function, then the induced risk measure preserves stochastic dominance between return distributions.

Example A.3.4 (Conditional Value-at-Risk)

Setting $\beta(\tau) = \min(\tau/\eta, 1)$ recovers CVaR_η , which averages the worst η -fraction of outcomes and yields a risk-averse policy.

Remark A.3.5

In practice the quantile samples $\tau \sim U([0, 1])$ are drawn independently for the online and target networks, which decorrelates the temporal-difference error.